

## ASSIGNMENT 2: ADDING SYSTEM CALLS

**Aim:** To add new system calls to a given kernel version and observe the results

**Requirement:** All work should be done on linux kernel version 4.19.210.

### ADDING KERNEL 4.19.210 TO SYSTEM

1. Created an azure student account using iiit mail id.
2. Created vm with **UBUNTU 16**, 4 core (linux-4.19.210).
3. Deployed vm on azure and used the public ip address to run the vm on the local machine.(using SSH client)
4. Created new folder in sudo mode and installed the pre-requisites required to run the kernel.
5. Go to above directory, created a new folder which contains all the required .c files, and then built and ran `sudo make` and `sudo make modules\_install`.
6. Made some configurations and rebooted the vm.

### ADDING A SYSTEM CALL TO KERNEL

1. Created a new directory in /linux-4.19.210 for a system call. In this text, let it be 'ass2'.
2. Created a .c files (eg q1.c) and add the necessary kernel space c code to it. We can use either `asm linkage long functionname(void)` or the macro `SYSCALL_DEFINE#(syscall_name)` where # is the number of arguments expected and `syscall_name` is the name given to it.
3. Create a file named 'Makefile' and add the line `'obj-y := q1.o'` to it, if the name of the c file is q1.c.
4. Navigate to / linux-4.19.210/include/linux and open `syscalls.h`
5. Add the definition of the function created in q1.c
6. Navigate to / linux-4.19.210/ and open Makefile
7. Add the name of the folder as 'ass2/' to core-y assignments as shown in subsequent headers.
8. Navigate to /linux-4.19.210/arch/x86/entry/syscalls and open `syscall_64.tbl` (or `syscall_32.tbl` if 32 bit system) and add appropriate entries to it. If function was made directly, we can make entry as shown for number 548, but if `SYSCALL_DEFINE#` macro was used, the following 3 lines should be used as references (x64(or32) \_\_sys\_ prefix should be added to the last column).
9. Finally, navigate to /linux-4.19.210 and run the following commands:

```
`sudo make modules_install`  
`sudo make install`
```

10. Updated the GRUB menu and restarted the vm and, then enter the same kernel again
11. Test the system call via a c code.

## SYSCALLS TO CREATE:

Four system calls are to be created, given below:

1. Printing a welcome message to kernel logs
2. Printing a given string to kernel logs
3. Printing current and parent process id to kernel logs
4. Recreating an existing syscall and implementing that (in this case, getpid() has been recreated)

## SYSCALL 64.TBL, MAKEFILE AND SYSCALLS.H

This list contains syscall entries and reference numbers for them. For each new syscall that we create, we need to add the corresponding entry. For the above 4 tasks, the last 4 entries in the given image were added (548, 549, 550, 551)

```
root@myVM: ~/linux-4.19.210/arch/x86/entry/syscalls
GNU nano 2.5.3 File: syscall_64.tbl

534      x32      preadv      __x32_compat_sys_preadv64
535      x32      pwritev     __x32_compat_sys_pwritev64
536      x32      rt_tgsigqueueinfo __x32_compat_sys_rt_tgsigqueueinfo
537      x32      recvmmsg    __x32_compat_sys_recvmmsg
538      x32      sendmmsg    __x32_compat_sys_sendmmsg
539      x32      process_vm_readv __x32_compat_sys_process_vm_readv
540      x32      process_vm_writev __x32_compat_sys_process_vm_writev
541      x32      setsockopt   __x32_compat_sys_setsockopt
542      x32      getsockopt   __x32_compat_sys_getsockopt
543      x32      io_setup     __x32_compat_sys_io_setup
544      x32      io_submit    __x32_compat_sys_io_submit
545      x32      execveat     __x32_compat_sys_execveat/ptregs
546      x32      preadv2      __x32_compat_sys_preadv64v2
547      x32      pwritev2     __x32_compat_sys_pwritev64v2
548      64      saihello      __x64_sys_saihelli
549      64      saiprint      __x64_sys_saiprint
550      64      saiprocess    __x64_sys_saiprocess
551      64      saigetpid     __x64_sys_saigetpid
```

Each system call was added in a different folder under the linux-4.19.210 directory, so each of the folders have to be added in the Makefile for the kernel itself (cust1 through 4 for each question 1 to 4)

Finally, the required definitions are also added into syscalls.h



```
Select root@myVM: ~/linux-4.19.210/include/linux
GNU nano 2.5.3 Fi

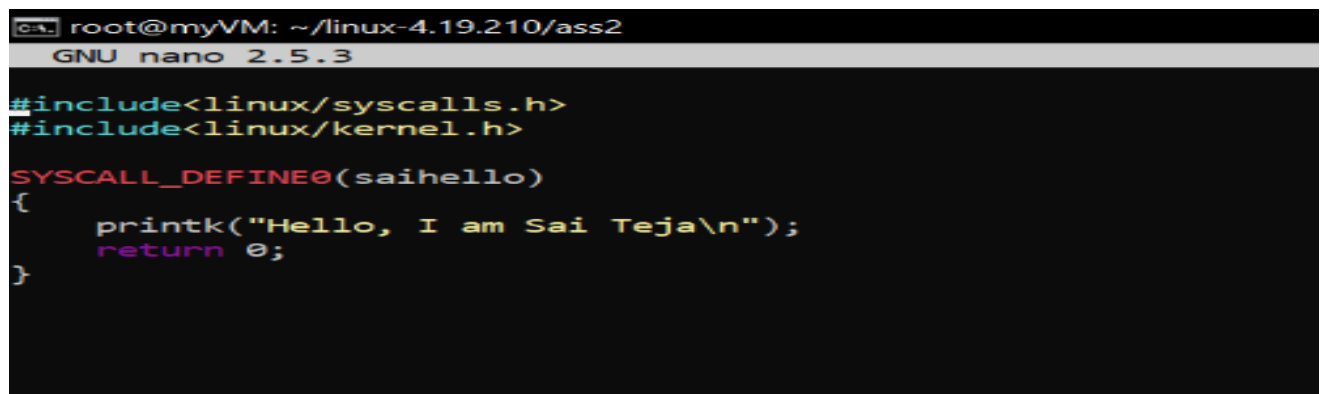
    return old;
}

asmlinkage long sys_saihello(void);
asmlinkage long sys_saiprint(char __user *);
asmlinkage long sys_saiprocess(void);
asmlinkage long sys_saigetpid(void);

#endif
_
```

## WELCOME MESSAGE TO KERNEL LOGS

1. The following code was added to saihello.c in /linux-4.19.210/ass2, along with the following Makefile containing 'obj-y := saihello.o'.



```
root@myVM: ~/linux-4.19.210/ass2
GNU nano 2.5.3

#include<linux/syscalls.h>
#include<linux/kernel.h>

SYSCALL_DEFINE0(saihello)
{
    printk("Hello, I am Sai Teja\n");
    return 0;
}
```

2. A new syscall entry was added in syscall\_64.tbl

## PRINTING SOME STRING FROM USER SPACE TO KERNEL SPACE

1. The following code was added to saiprint.c in /linux-4.19.210/ass2, along with the following Makefile containing 'obj-y := saiprint.o'.

```
root@myVM: ~/linux-4.19.210/ass2
GNU nano 2.5.3

#include <linux/kernel.h>
#include <linux/linkage.h>
#include <linux/syscalls.h>
#include <linux/uaccess.h>

SYSCALL_DEFINE2(saiprint,
                char __user *, src,
                int, len)
{
    char buf[256];
    unsigned long lenleft = len;
    unsigned long chunklen = sizeof(buf);
    while( lenleft > 0 ){
        if( lenleft < chunklen ) chunklen = lenleft;
        if( copy_from_user(buf, src, chunklen) ){
            return -EFAULT;
        }
        lenleft -= chunklen;
    }
    printk("%s\n", buf);
    return 0;
}
```

Here, SYSCALL\_DEFINE2 is a macro that is creating a function `asmlinkage long saiprint(char __user * str,int len)`, with one argument `str` and one `int`. The '`__user`' marks the pointer as one to the user space. '`Copy_from_user`' is an API call that allows copying data from a pointer in userspace to kernelspace buffer. If there is some sort of error while fetching information, `EFAULT` error flag will be returned.

2. A new syscall entry is added in `syscall_64.tbl`. It is different from the first entry in that it has a new `sys_` prefix and `__x64_sys_` prefix for alias and syscall name respectively. This is a result of using the `SYSCALL_DEFINE#` macro.

## PRINTING CURRENT AND PARENT PROCESS ID

1. The following code was added to `saiprocess.c` in `/linux-4.19.210/ass2`, along with the following Makefile containing '`obj-y := saiprocess.o`'.

```
Select root@myVM: ~/linux-4.19.210/ass2
GNU nano 2.5.3 File

#include<linux/syscalls.h>
#include<linux/kernel.h>
#include<linux/cred.h>
#include<linux/sched.h>

SYSCALL_DEFINE0(saiprocess)
{
    struct task_struct *parent=current->parent;
    printk("parent_process_pid: %d \n", parent->pid);

    printk("current_process_pid: %d \n", current->pid);
    return 0;
}
```

'Current' refers to a structure of type 'task' that stores information about the current process (that includes the entirety of the Process Control Block). The current->pid element will have the current process id. The parent->pid call will return the pid of the parent process given a task structure for current process.

2. Appropriate entry in syscall\_64.tbl is added

#### RECREATING A SYSCALL: 'getpid()'

1. The following code was added to saigetpid.c in /linux-4.19.210/ass2, along with the following Makefile containing 'obj-y := saigetpid.o'.

```
root@myVM: ~/linux-4.19.210/ass2
GNU nano 2.5.3 File

#include<linux/syscalls.h>
#include<linux/kernel.h>
#include<linux/cred.h>
#include<linux/sched.h>

SYSCALL_DEFINE0(saigetpid)
{
    printk("%u\n",task_tgid_vnr(current));
    return task_tgid_vnr(current);
}
```

The uidgid.h header contains the k\_uid structure that will have the user id stored in it. Syscall getuid() itself fetches this information from this place. Current\_uid() will return the k\_uid structure and val will return the uid unsigned long from it.

2. Appropriate entry in syscall\_64.tbl is added

## OUTPUTS

```
root@myVM: ~
GNU nano 2.5.3 File: 1

#include<stdio.h>
#include<unistd.h>
#include<sys/syscall.h>

int main(int argc, char** argv){
    syscall(548);
    syscall(549, "Hello World ", 100);
    syscall(550);
    printf("%ld", syscall(551));
    return 0;
}
```

The above code has been used to test the newly created syscalls 548 to 551.

The kernel logs are as follows

```
root@myVM: ~
root@myVM:~# ./a.out
2062root@myVM:~# dmesg
[ 422.979801] Hello, I am Sai Teja
[ 422.979806] Hello World
[ 422.979807] parent_process_pid: 1998
[ 422.979808] current_process_pid: 2062
[ 422.979809] 2062
root@myVM:~#
```

Parent and current process ids are different. Why?

The process id refers to the process that is handling both the syscall and the function that called the syscall. That is, in this case saiprocess.c and test.c are part of the same process. Calling another function does not necessarily create a new process, hence 2062 in this case refers to our currently running functions.

The parent process id 1998 refers to the parent process of test.c, the code within which the system call is being tested.